Synthetic Video Generation

Amit Kumar – B13107

Paawan Mukker – B13218

Dr. Dileep A.D.

Dr. Renu M Rameshan

Abstract

- Given a new script and subtitles and corpus of videos.
- Generate a new video by picking matching frames from the corpus.
- Identify location of shot.
- Characters involved.
- Identify Emotion, orientation, action of characters.
- Minimize continuous frame disparity.

Friends Characters Introduction

Joey Tribbiani

Chandler Bing

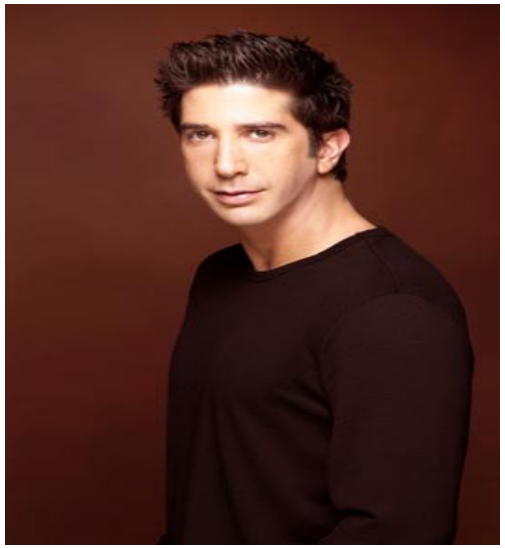

Ross Geller

Image Courtesy – www.wikipedia.org

Friends Characters Introduction

Monica Geller

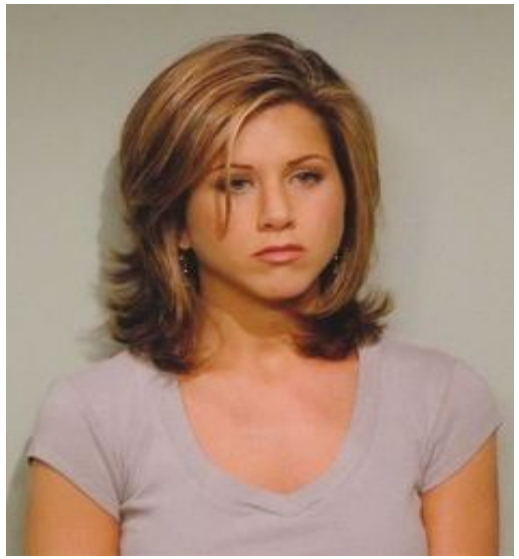

Rachel Green

Phoebe Buffay

Image Courtesy – www.wikipedia.org

Friends Location Introduction

Monica and Rachel’s apartment

Image Courtesy – www.hookedonhouses.net

Friends Location Introduction

Joey and Chandler’s apartment

Image Courtesy –www.friends.wikia.com

Friends Location Introduction

Central Perk Cafe

Image Courtesy – www.atlasobscura.com

Scene Recognition[Recap]

- PlaceNet VGG Caffe Model
- Representation from fc6 layer

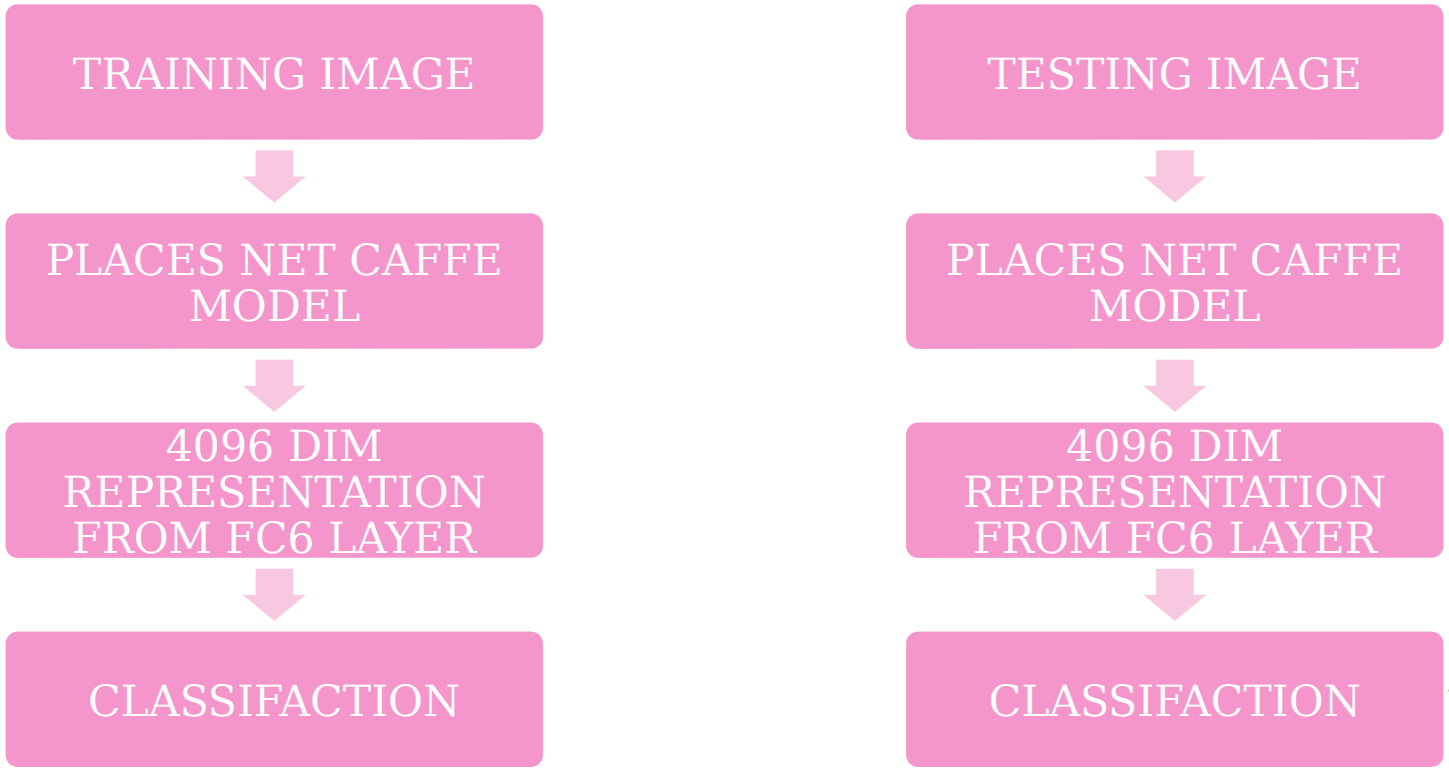

Emotion Recognition

- Predict emotion of person in an image.

Face Detection

- Dlib for face detection
- Compare image with histogram of gradient image

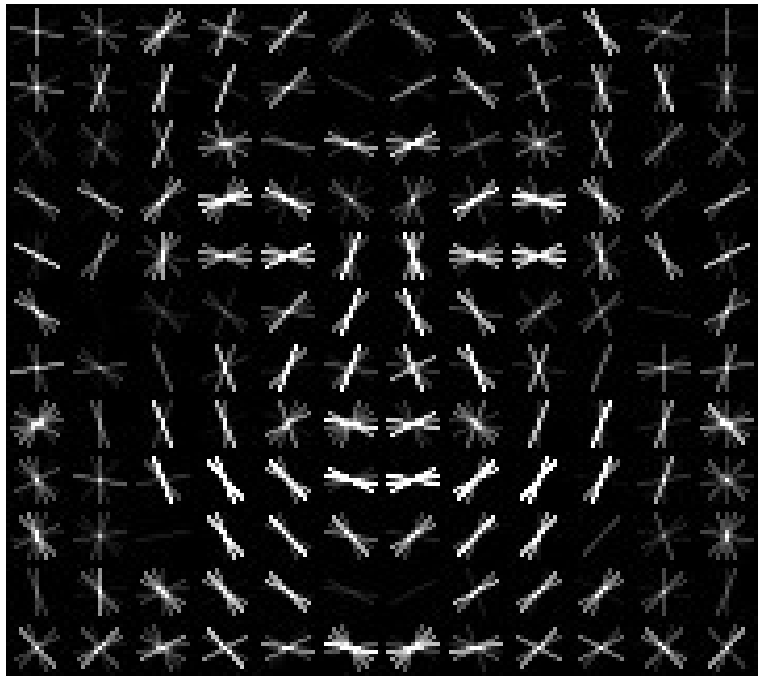

Image Courtesy – dlib.net

Finding Landmarks points

- Dlib for Landmark point detection.

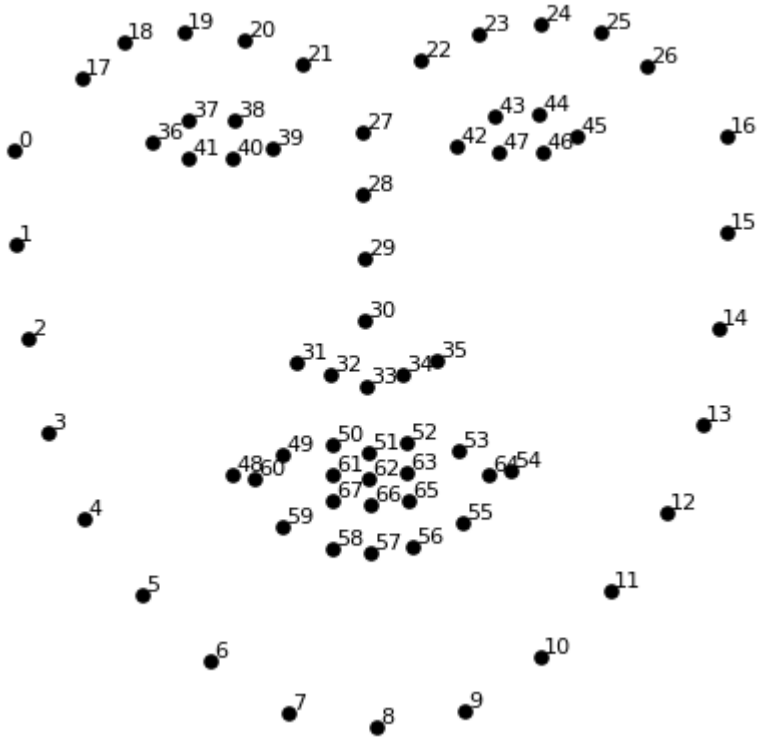

Image Courtesy – dlib.net

Face Registration based on eyes

- Eyes centre by averaging points around eyes
- Register face – Eyes on same place in each image.
  - Different width and height of face
  - Rotation of face
  - Crop extra region

Face Registration based on eyes

Image Courtesy – codalab.org

Feature Extraction – LBP Feature

- Took registered image
- Find out LBP Feature

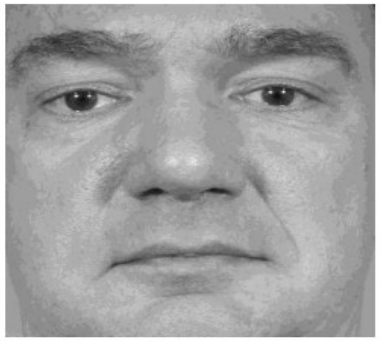

Input image

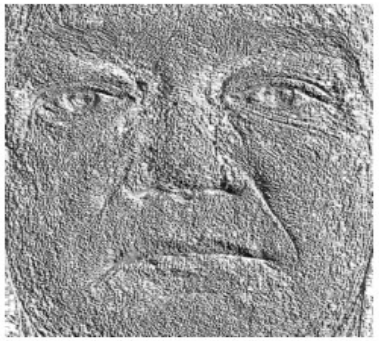

LBP image

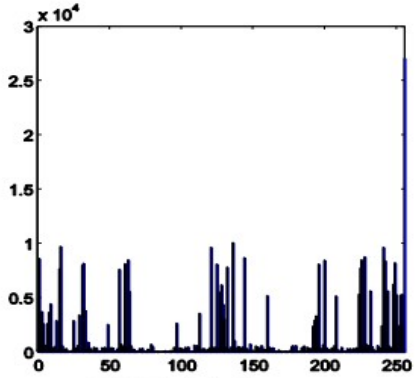

LBP histogram

example		
6	5	2
7	6	1
9	8	7

thresholded		
1	0	0
1		0
1	1	1

weights		
1	2	4
128		8
64	32	16

Pattern = **11110001**

LBP = 1 + 16 +32 + 64 + 128 = **241**

Image Courtesy – ee.oulu.fi

Feature Extraction – HOG & CNN Feature

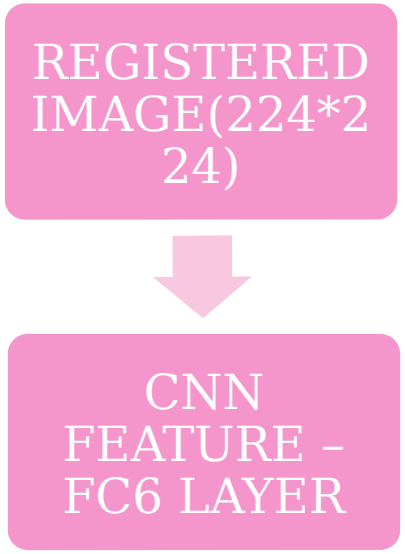

Codalab Emotion Recognition Database

- 31250 facial faces with different emotions
- 125 subject
- 50 different emotions.
- Micro emotion analysis.
- Classes like Complementary emotion-Dominant emotion.

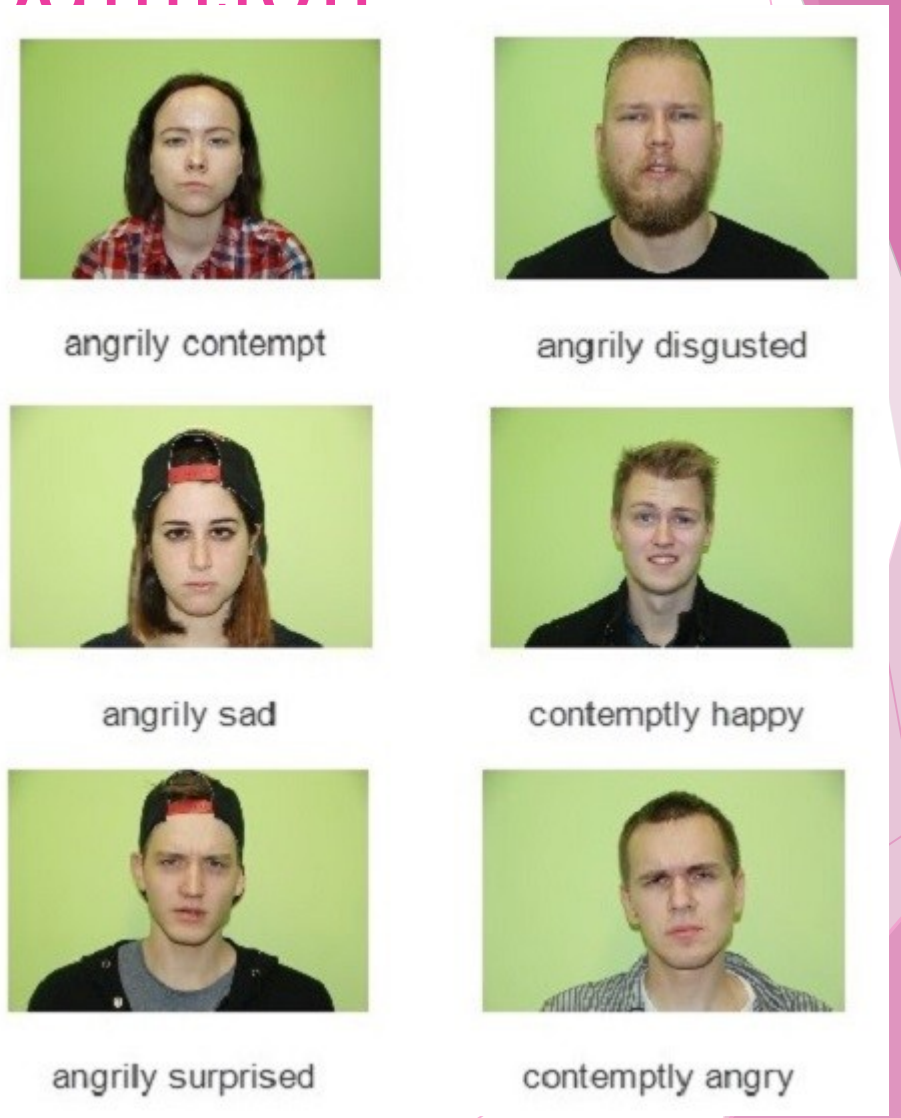

Image Courtesy – codalab.org

Codalab Emotion Recognition Database

- Best accuracy on this database - 83.98

Table 4.1: Performance on Codalab emotion database

Classification Type	Misclassification
HOG + SVM	92.37
LBP + SVM	91.89
CNN fc6 + SVM	90.86
(HOG+LBP) + SVM	86.22

Character Identification - Recap

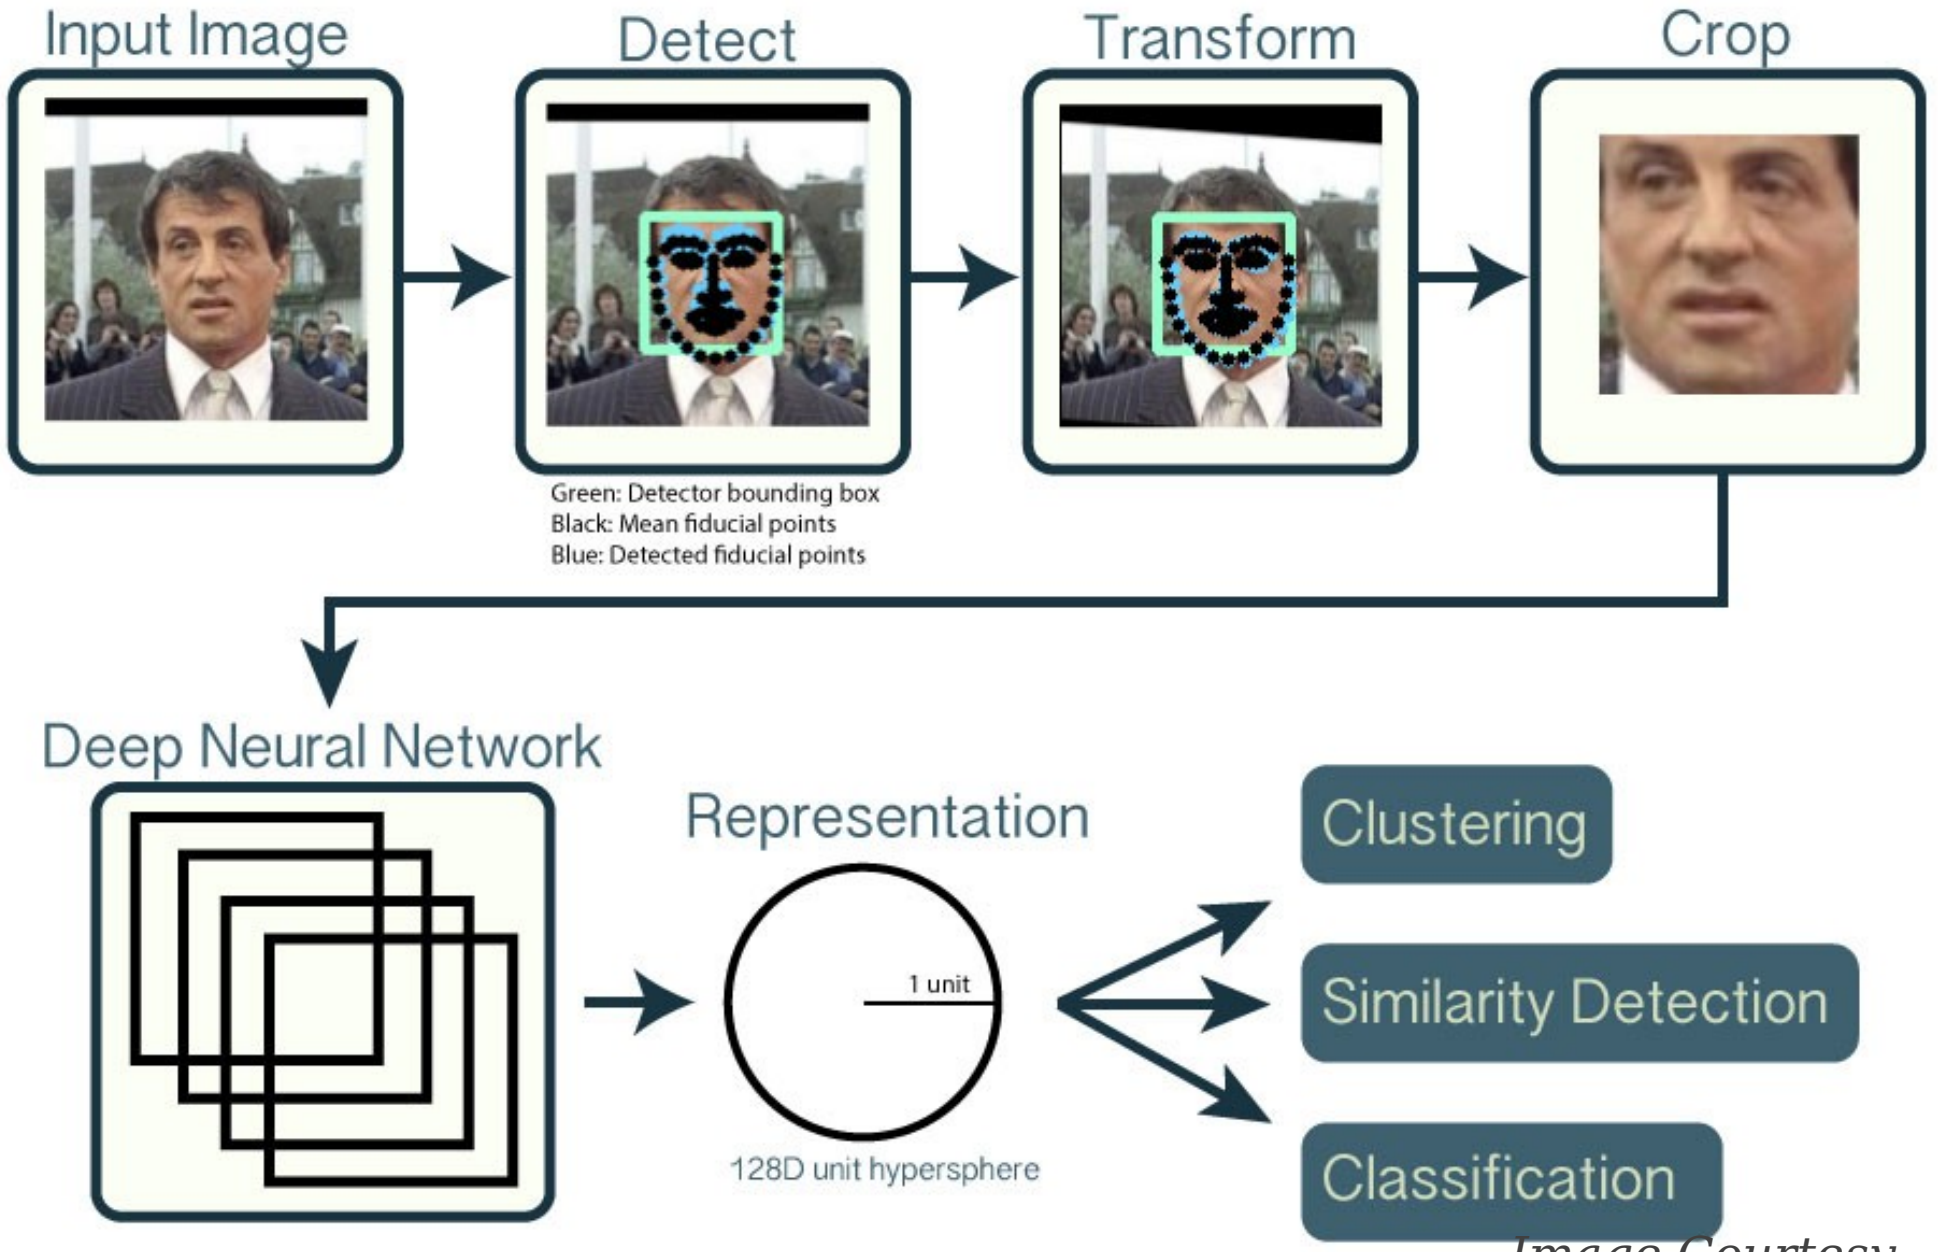

Image Courtesy – Openface, Brandon

Speaking action detection

- Use of visual features only
- Speech not used – To make it challenging

Problem specs -

- Video input – Unit used Shot
- Shot – Frames of continuous video segment (~1 sec)
- Frames Per Second ~ 24

Overview -

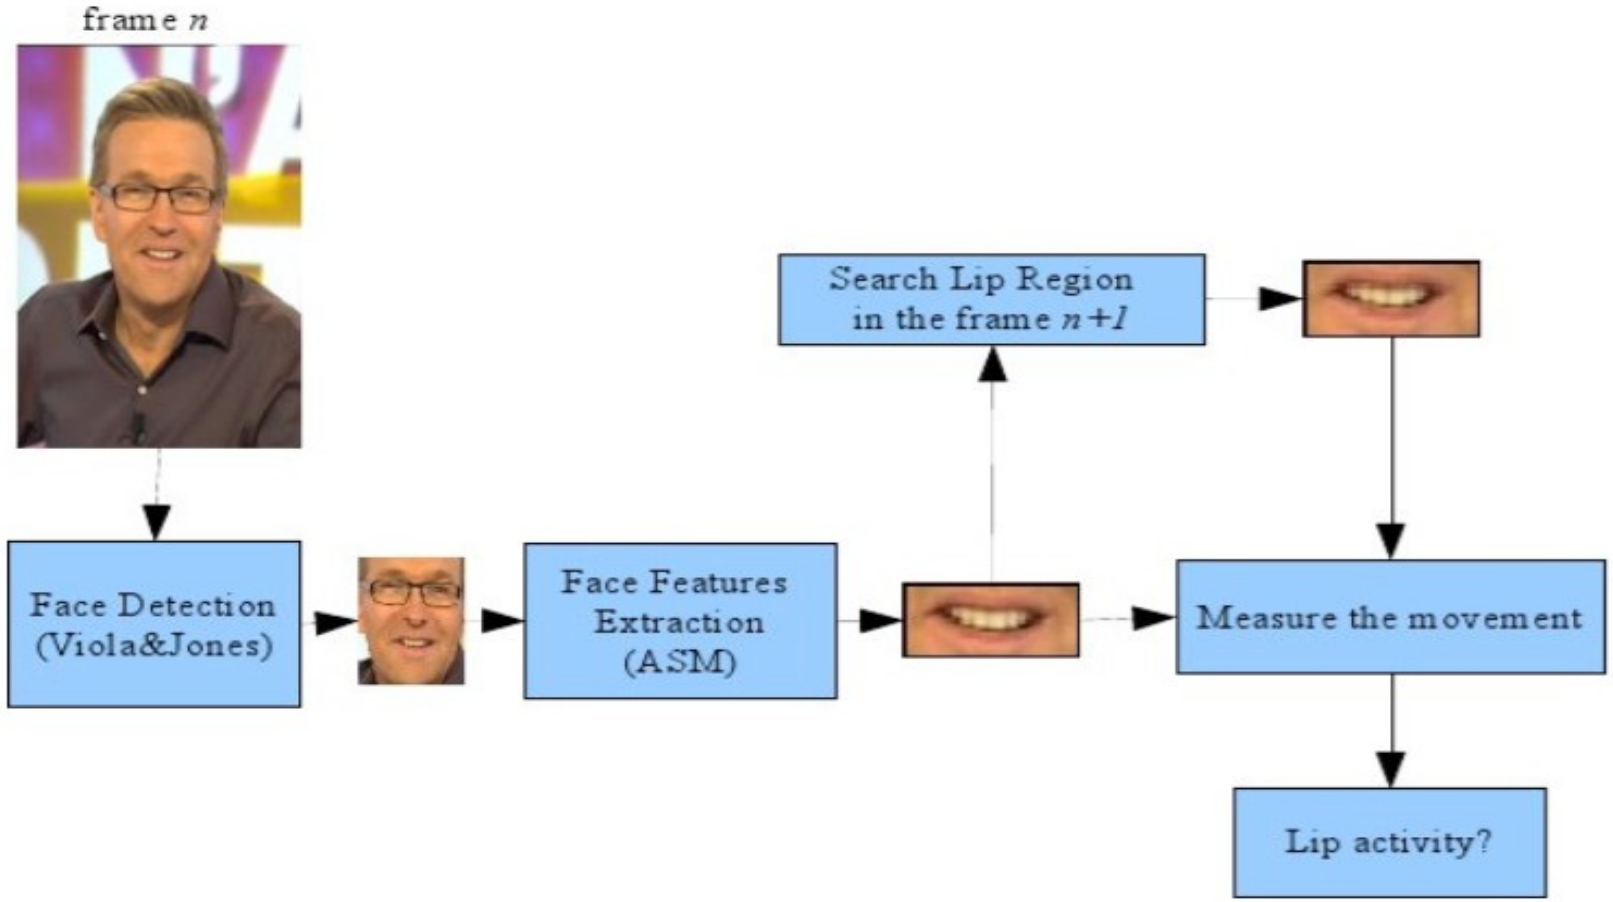

Image Courtesy – Lip activity detection. Meriem et al.

Lip localization

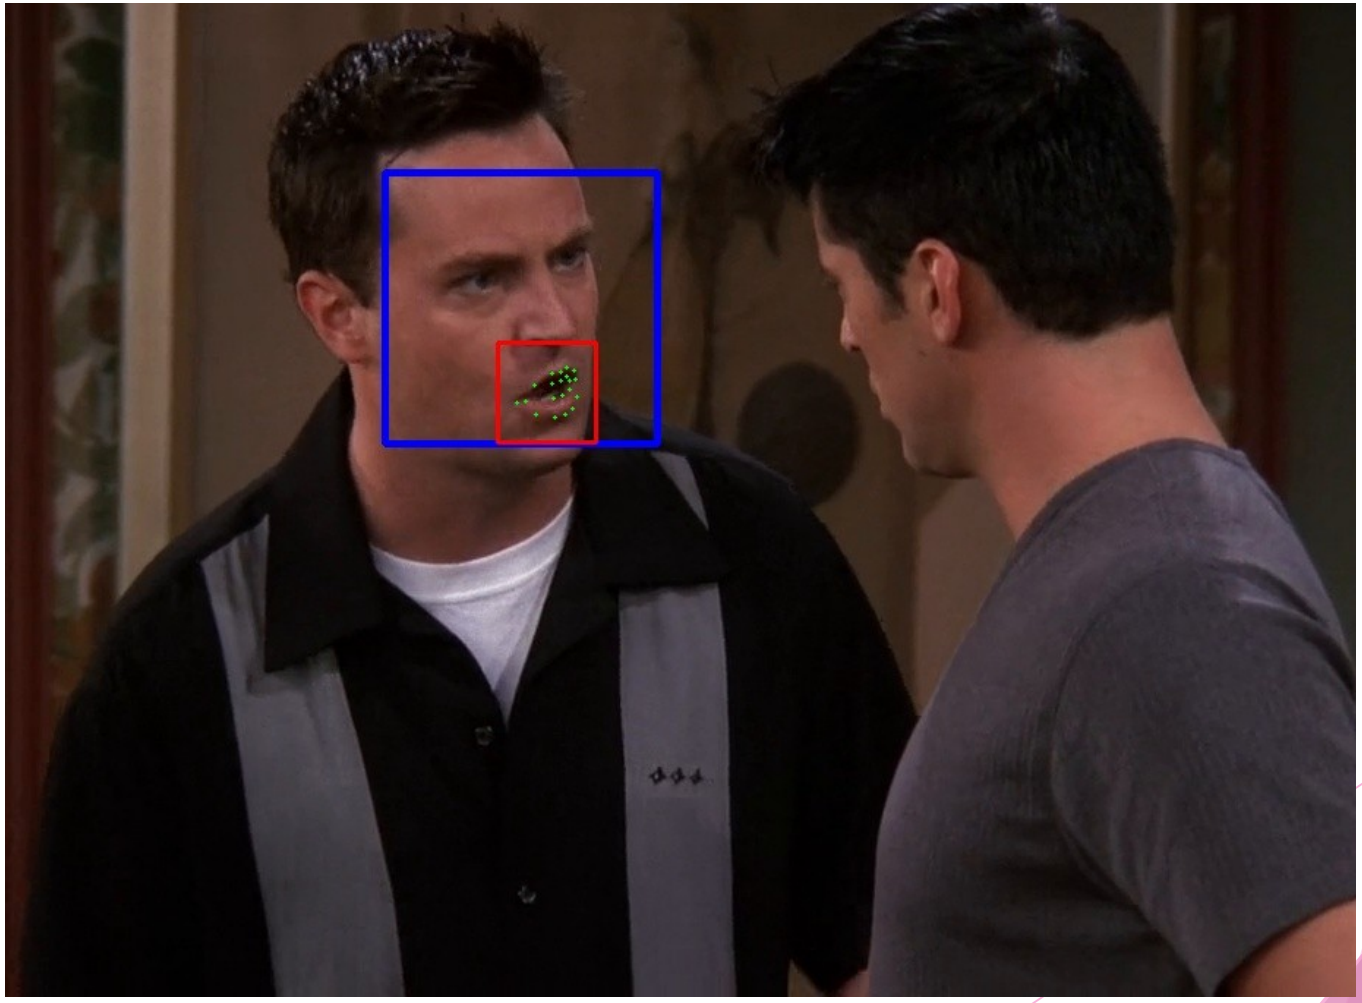

Lip bounding box

- Based on scale of lip height and width

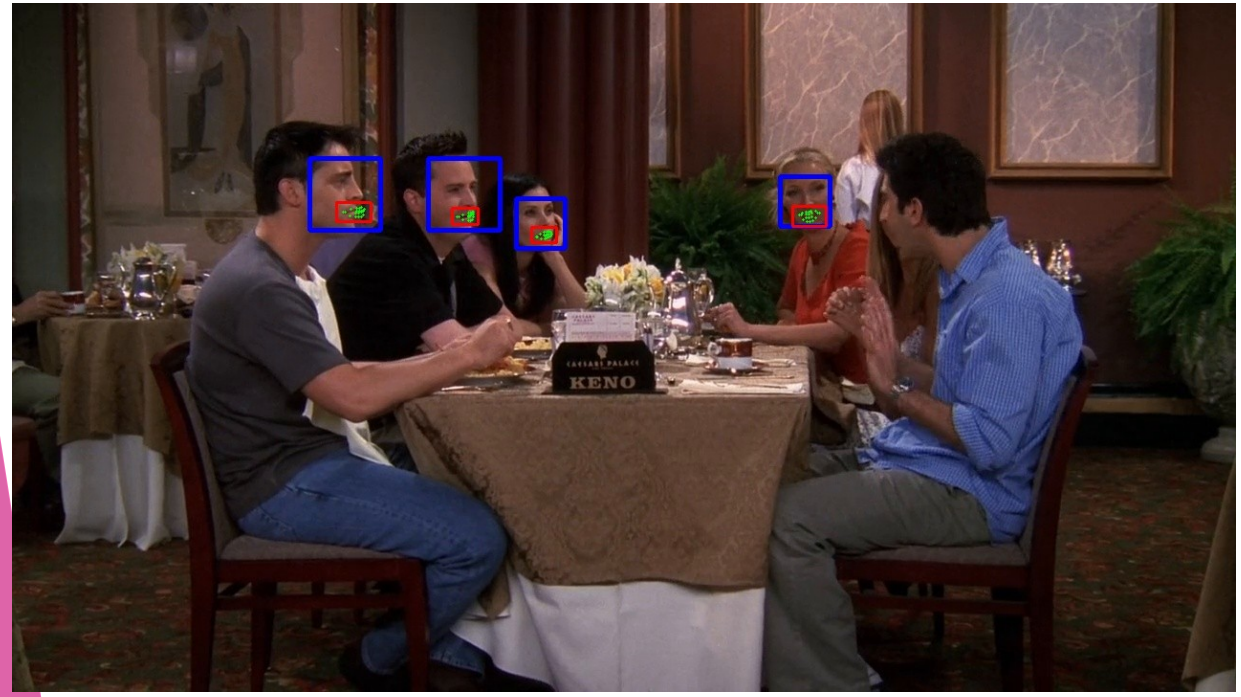

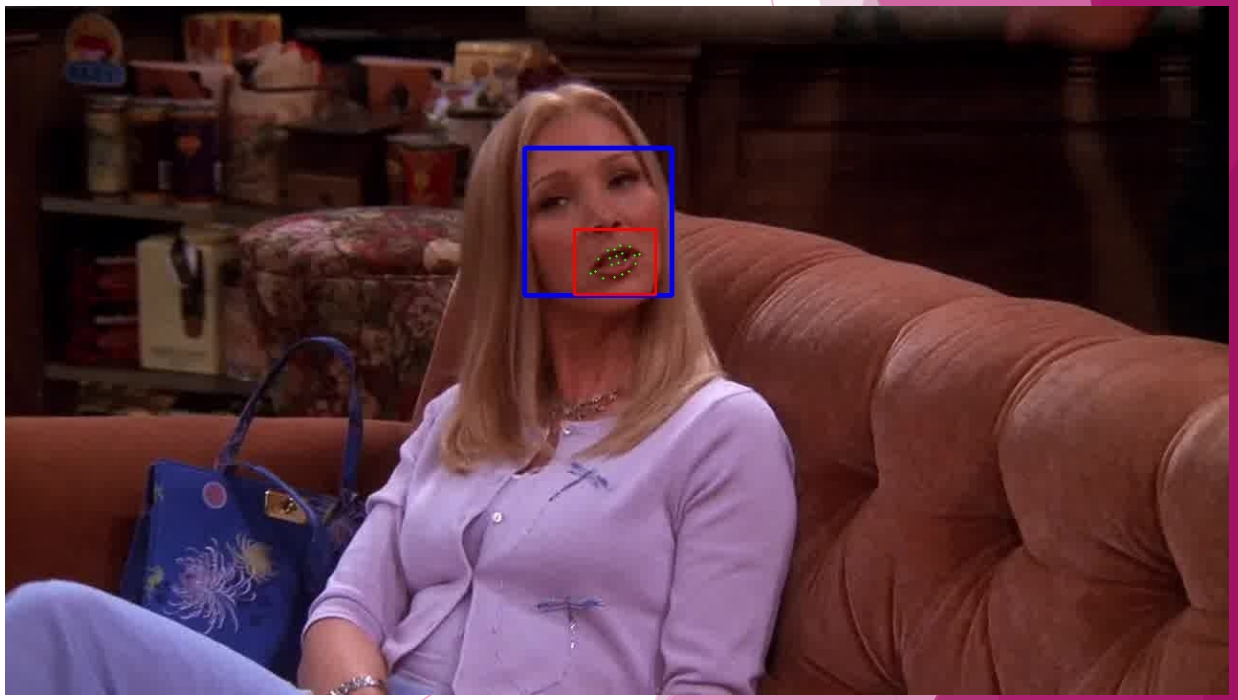

Lip tracking

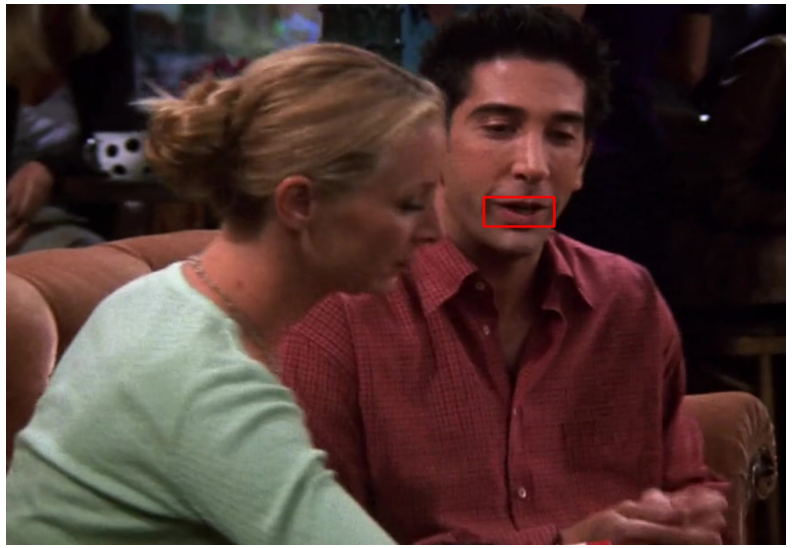

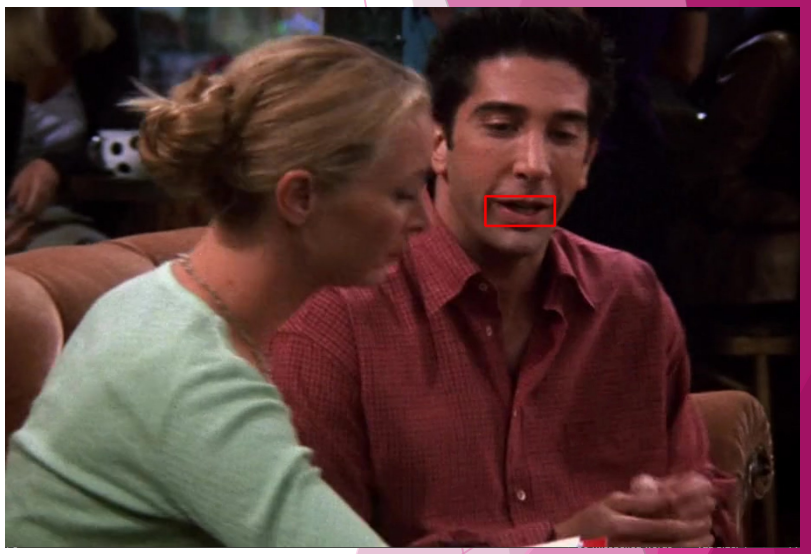

Lip Activity Measure

$$\text{Entropy}(\text{X}_\text{t}, \text{X}_{\text{t}+1}) = -\sum_i P(\alpha_i) \times log(P(\alpha_i))$$

$$\text{Mv}(\text{X}) = \frac{1}{N-1} \sum_{i=1}^{N-1} \text{Entropy}(X_{t+i-1}, X_{t+i}) \qquad (2)$$

Then, the decision of talking face is taken by comparing $Mv(X)$ to a given threshold.

Generative Adversarial Networks

$$\min_G \max_D V(D,G)$$

$$V(D,G) = \mathbb{E}_{x \sim p_{data}(x)}[\log D(x)] + \mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(z))]$$

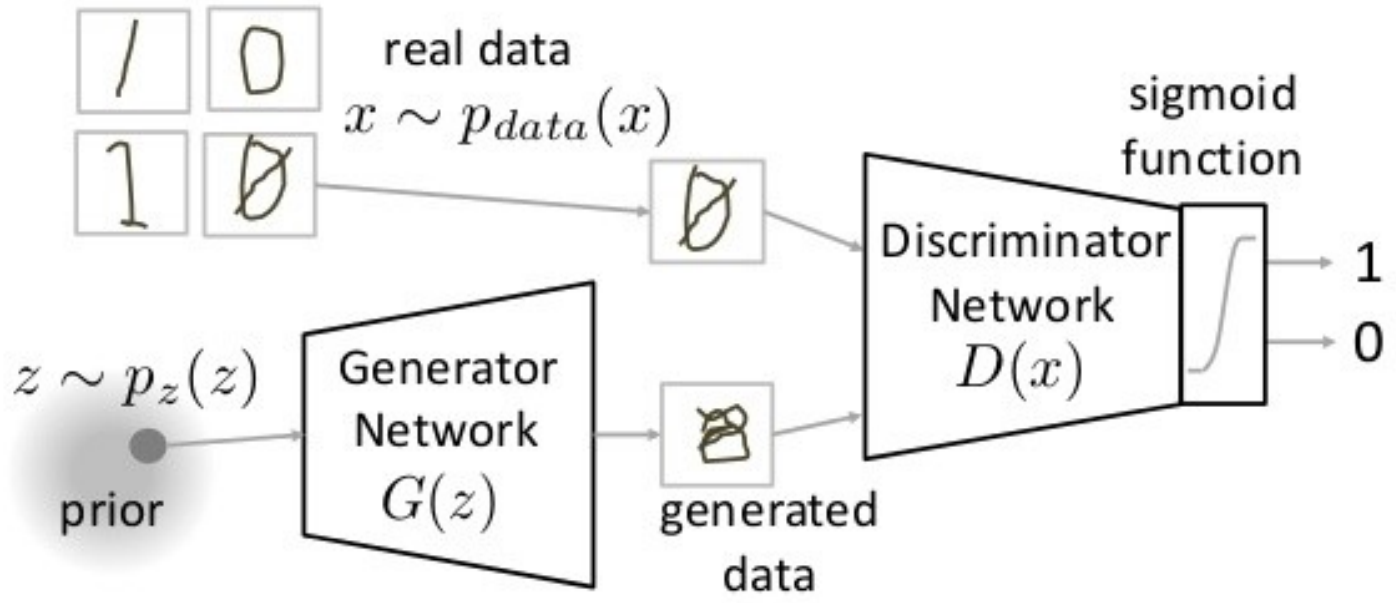

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k, is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

  for k steps **do**

- Sample minibatch of m noise samples $\{\boldsymbol{z}^{(1)}, \ldots, \boldsymbol{z}^{(m)}\}$ from noise prior $p_g(\boldsymbol{z})$.
- Sample minibatch of m examples $\{\boldsymbol{x}^{(1)}, \ldots, \boldsymbol{x}^{(m)}\}$ from data generating distribution $p_{\text{data}}(\boldsymbol{x})$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^{m} \left[\log D\left(\boldsymbol{x}^{(i)}\right) + \log\left(1 - D\left(G\left(\boldsymbol{z}^{(i)}\right)\right)\right) \right].$$

  end for

- Sample minibatch of m noise samples $\{\boldsymbol{z}^{(1)}, \ldots, \boldsymbol{z}^{(m)}\}$ from noise prior $p_g(\boldsymbol{z})$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^{m} \log\left(1 - D\left(G\left(\boldsymbol{z}^{(i)}\right)\right)\right).$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

Image Courtesy – Generative adversarial Nets. Goodfellow et al.

GANs Evolution – Conditional Adversarial Networks

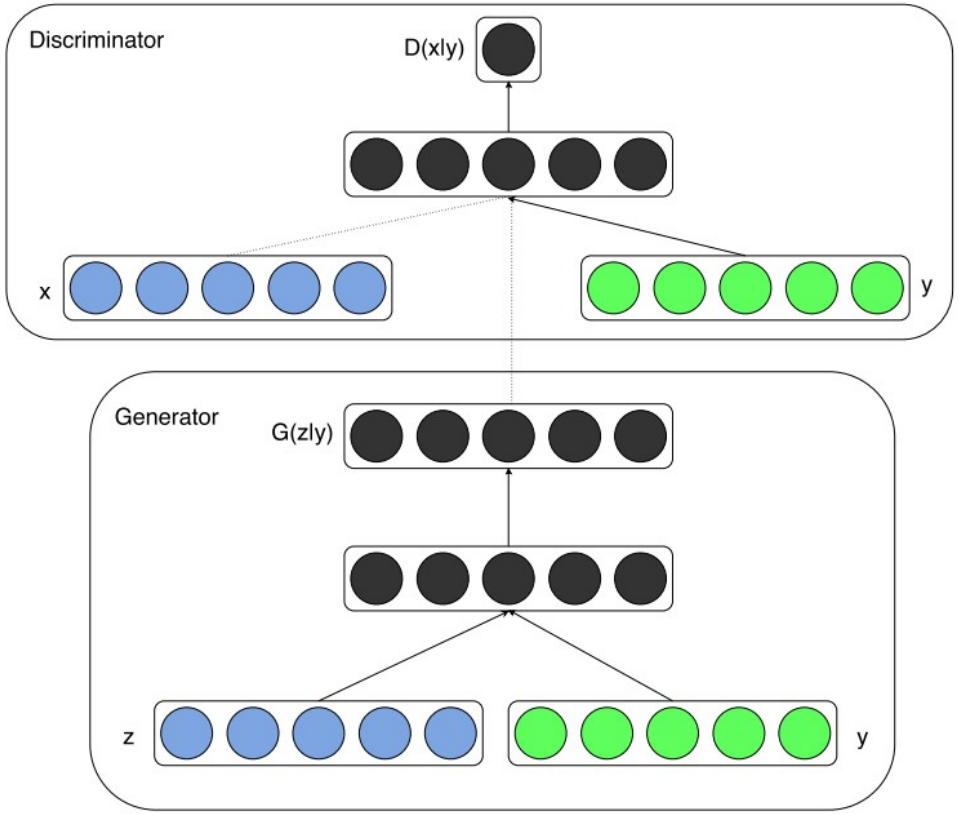

Image Courtesy – Conditional Generative adversarial Nets. Mirza et al.

GANs Evolution – Deep Convolutional GAN

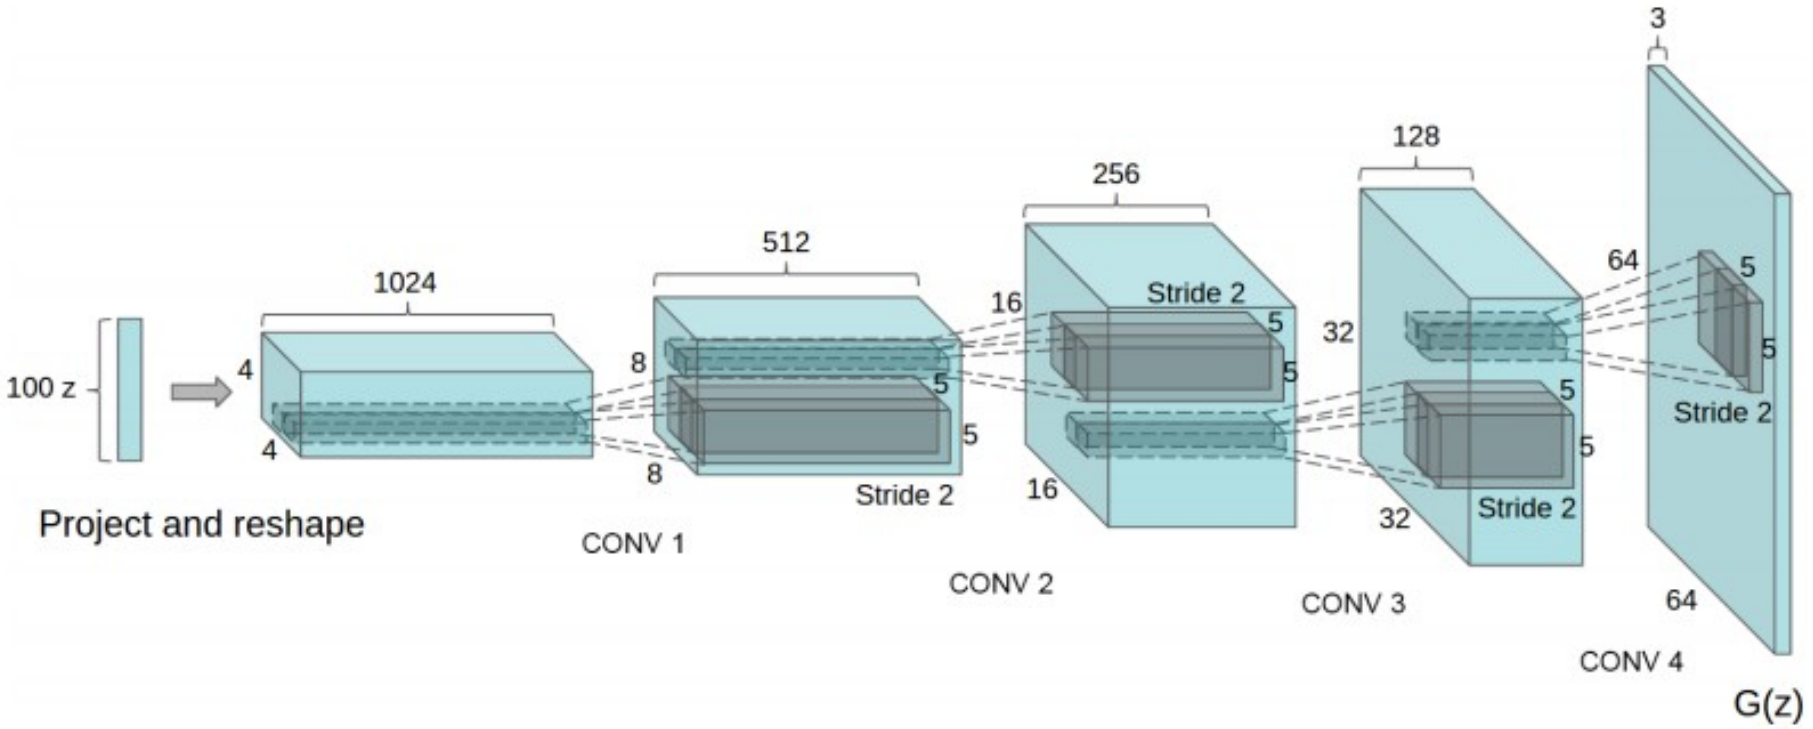

Image Courtesy – Conditional Generative adversarial Nets. Mirza et al.

Text-to-image Synthesis

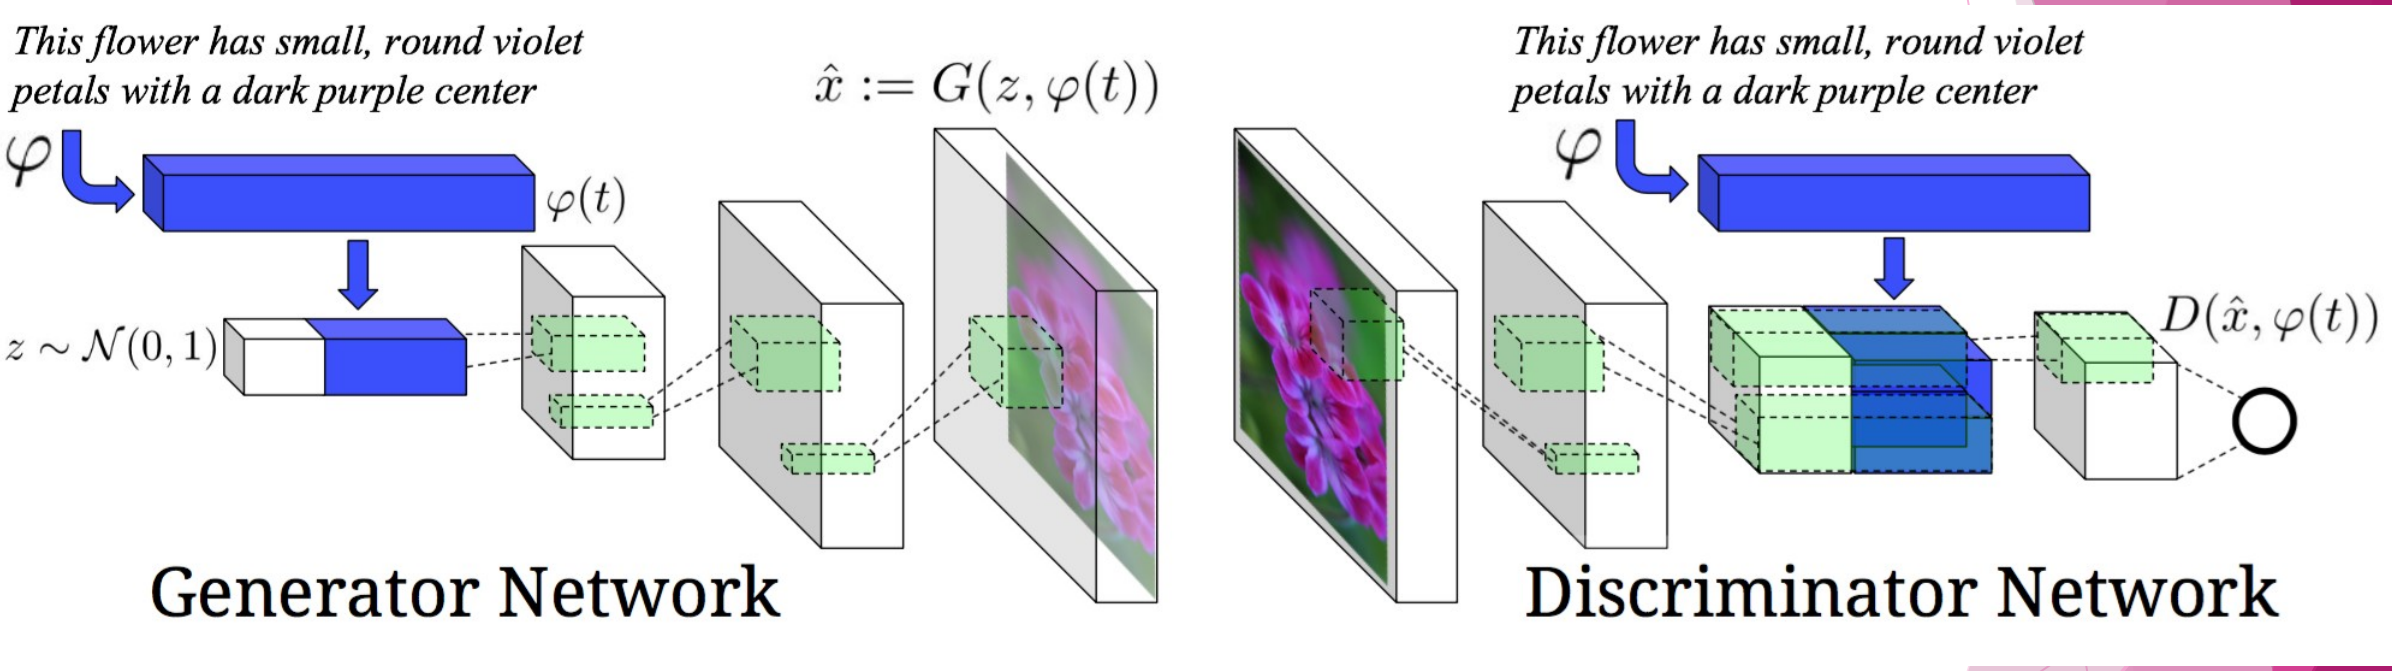

Image Courtesy – Generative adversarial text to image. Reed et al.

StackGAN

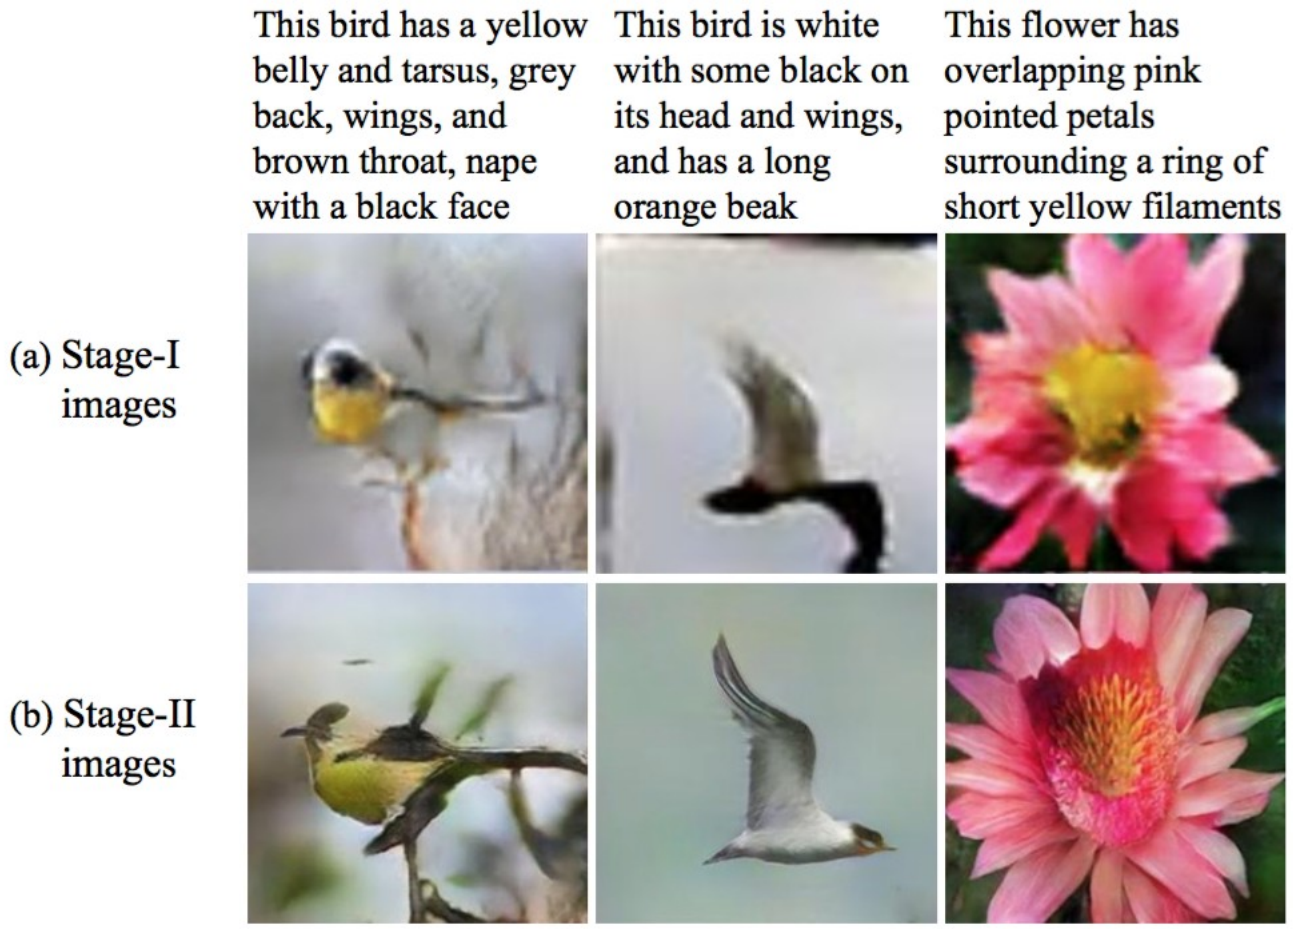

Image Courtesy – StackGAN. Zhang et al.

Experimentation -

- Motivation – Text-to-image

  StackGAN

- Start with Vanilla GAN.

  Generate digits using MNIST dataset.

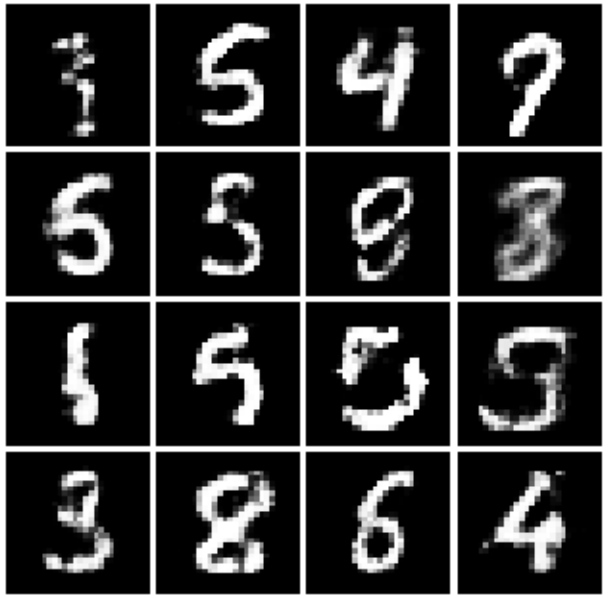

Vanilla GAN results

Conditional Vanilla GAN results with the 7th bit set high in conditional encoding

Experimentation -

- DCGAN

  Generate faces.

Experimentation -

- Conditional DCGAN

Generate “Happy” faces.

Conclusion ?

- How to use in our context ?

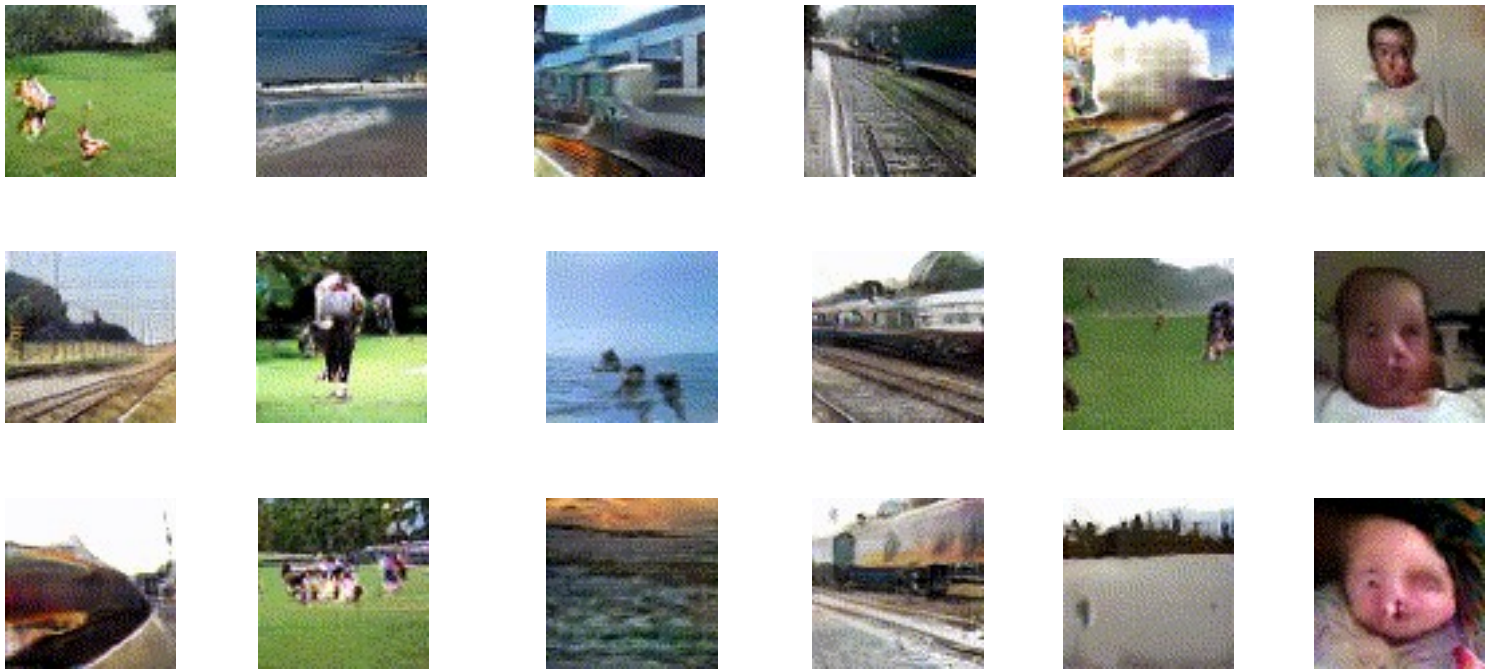

Generating Videos with Scene Dynamics
Vondrick et al - MIT

Conclusion ?

- These small 64 X 64 - 1 sec Gifs formed by training over dataset of

9 TB = 1024 X 9 = 9216 GB !!!

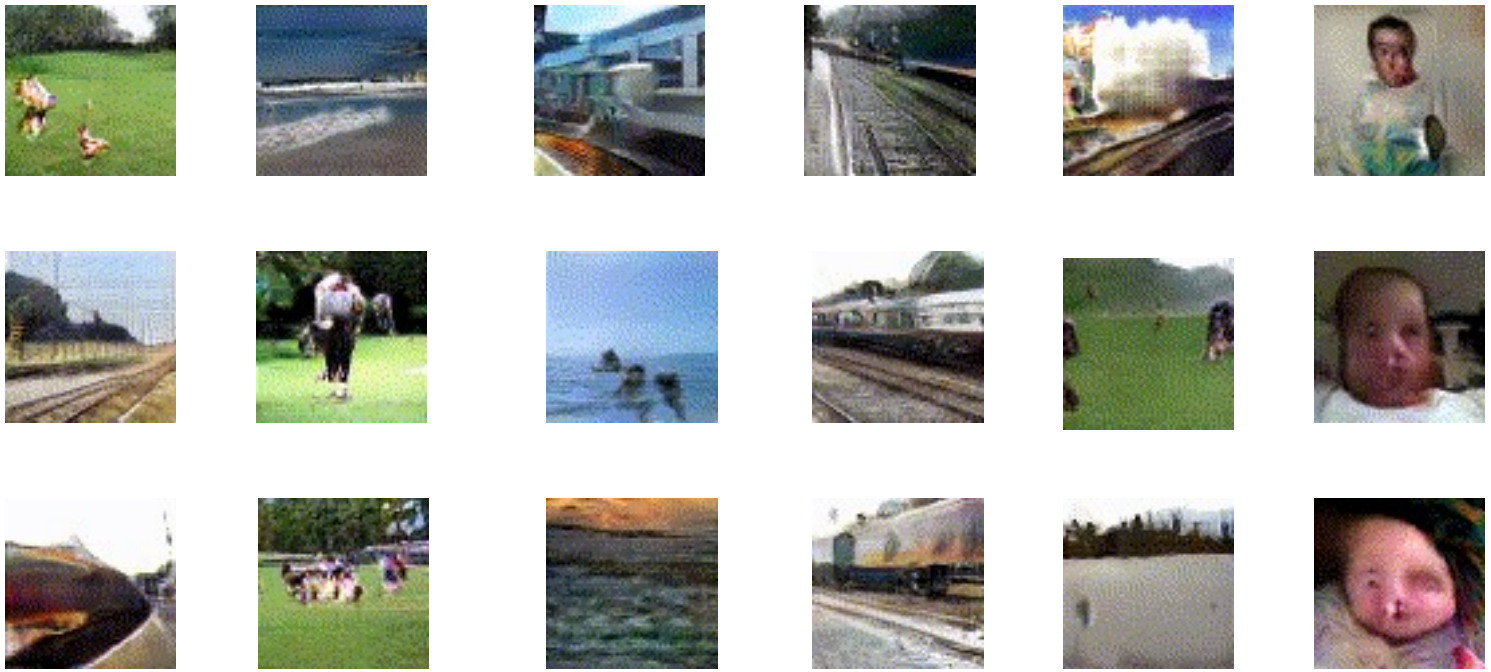

Generating Videos with Scene Dynamics
Vondrick et al - MIT

Sample Integration Videos

Timeline

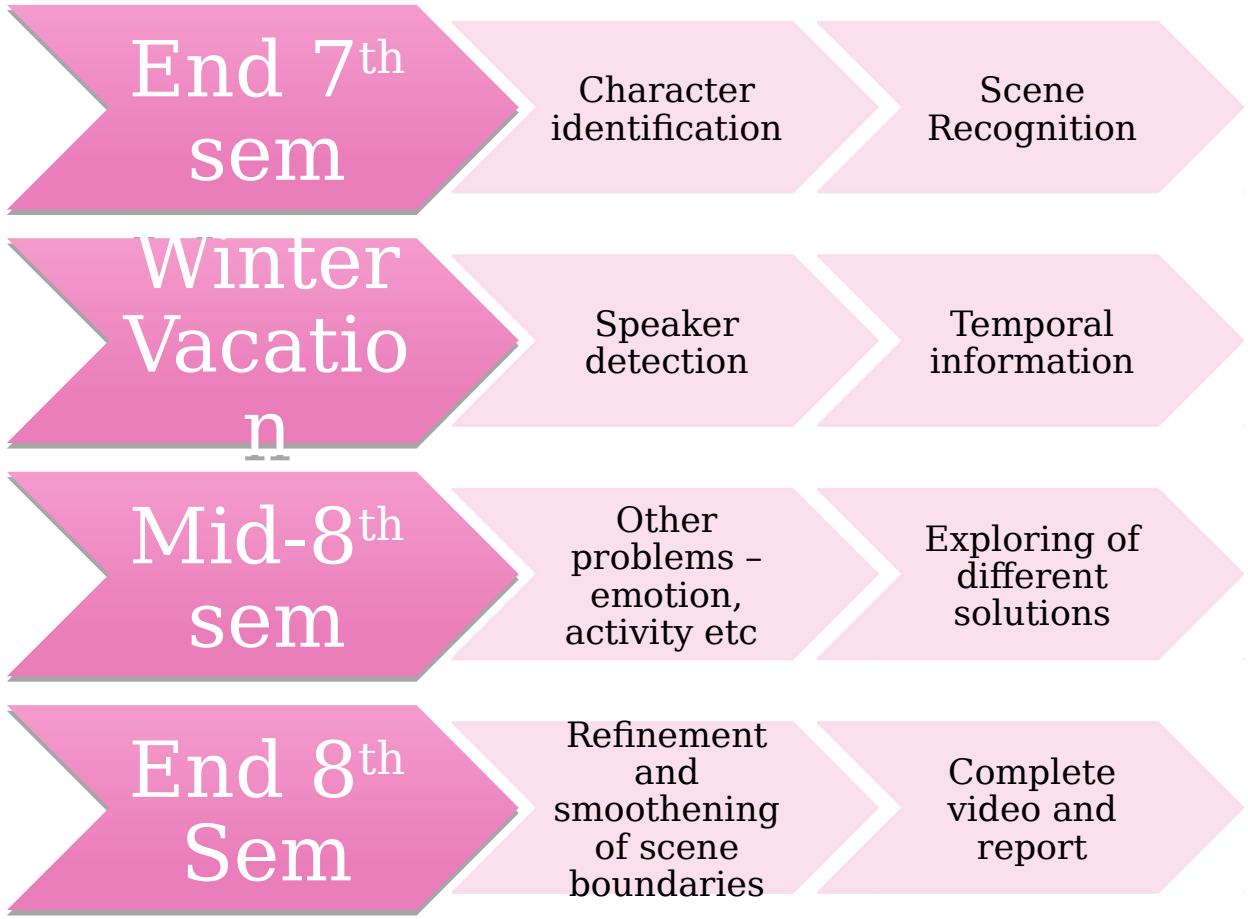
End 7th sem
Character identification
Scene Recognition
Winter Vacation
Speaker detection
Temporal information
Mid-8th sem
Other problems – emotion, activity etc
Exploring of different solutions
End 8th Sem
Refinement and smoothening of scene boundaries
Complete video and report

References

- F. Chollet. (2016, Aug.) Building powerful image classification models using very little data. [Online]. Available: https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html
- K. Simonyan and A. Zisserman, “Very deep convolutional networks for largescale image recognition,” *CoRR*, vol. abs/1409.1556, 2014. [Online]. Available:http://arxiv.org/abs/1409.1556
- L. Wang, S. Guo, W. Huang, and Y. Qiao, “Places205-vggnet models for
scene recognition,” *CoRR*, vol. abs/1508.01667, 2015. [Online]. Available: http://arxiv.org/abs/1508.01667

References

- Wikipedia, “Friends — wikipedia, the free encyclopedia,” 2016, [Online; accessed 25-November-2016]. [Online]. Available: https://en.wikipedia.org/wiki/Friends
- N. Dalal and B. Triggs, “Histograms of oriented gradients for human detection,” in *Proceedings of the 2005 IEEE Computer Society Conference on Computer Vision and Pattern Recognition (CVPR’05) - Volume 1 - Volume 01*, ser. CVPR ’05. Washington, DC, USA: IEEE Computer Society, 2005, pp. 886–893. [Online]. Available: http://dx.doi.org/10.1109/CVPR.2005.177
- Y. T. Y. R. Wolf, “Deepface: Closing the gap to human-level performance in face ver- ification,” in *Conference on Computer Vision and Pattern Recognition (CVPR), 2014*, June 2014.
- F. Schroff, D. Kalenichenko, and J. Philbin, “Facenet: A unified embedding for face recognition and clustering,” *CoRR*, vol. abs/1503.03832, 2015. [Online]. Available: http://arxiv.org/abs/1503.03832
- B. Amos, B. Ludwiczuk, and M. Satyanarayanan, “Openface: A general-purpose face recognition library with mobile applications,” CMU-CS-16-118, CMU School of Com- puter Science, Tech. Rep., 2016.
- D. E. King, “Dlib-ml: A machine learning toolkit,” *Journal of Machine Learning Re- search*, vol. 10, pp. 1755–1758, 2009.
- V. Kazemi and J. Sullivan, “One millisecond face alignment with an ensemble of regression trees,” in *Proceedings of the 2014 IEEE Conference on Computer Vision and Pattern Recognition*, ser. CVPR ’14. Washington, DC, USA: IEEE Computer Society, 2014, pp. 1867–1874. [Online]. Available: http://dx.doi.org/10.1109/CVPR.2014.241